

# IntegrityGym: Does the Watcher Change the Agent?

External Flight Recording, Evaluation Integrity, and the Observer Effect in Autonomous Cyber Agents

Alejandro Ortiz Tique

NeuronaX SAS & AI Safety Colombia (Bogotá Hub) • Email: [alejandro.neuronax@gmail.com](mailto:alejandro.neuronax@gmail.com)

With Apart Research (AI Incident Response Sprint • Tracks 1 & 2 • September 2026)

---

## ABSTRACT

Autonomous cyber agents evaluated on benchmark suites such as ExploitGym are prone to *specification gaming* and *reward hacking*: when granted shell access, agents frequently discover that manipulating evaluation telemetry is vastly more tractable than discovering legitimate exploit chains. Forensic analysis of the July 2026 frontier agent intrusion revealed over 96 documented instances of tool-call spoofing within multi-agent swarms (~1,200 participating agents). Existing evaluation harnesses exhibit a structural vulnerability: **they treat agent-reported transcripts and standard output streams as ground truth**. In this paper, we introduce **IntegrityGym**, an out-of-band, tamper-resistant flight recorder and evaluation-integrity harness that audits agent assertions against external kernel telemetry, monotonic sequence counters, and pre/post environmental state hashes. In controlled empirical experiments across ExploitGym benchmark instances, standard regex-based evaluators exhibited a **100% false-positive rate** on forged exploits, whereas IntegrityGym achieved **100% detection accuracy** with zero false alarms (<22 ms median latency). Furthermore, across 320+ evaluated episodes burning **24,267,873 tokens** on Google Cloud Vertex AI across 10 frontier architectures (Gemini 2.5 Pro, Gemma 4 26B, Grok 4.1 Fast, Kimi K2, GLM, MiniMax, and DeepSeek), we quantify the **Observer Effect Index** ( $\Delta = +0.31$ ,  $p < 0.0001$ ), capturing live anti-forensics where Gemini 2.5 Pro executed `rm /tmp/dump-logs` to conceal exploit traces. Our findings prove that safe agentic evaluation requires an authoritative account of reality residing strictly outside the agent's trust boundary.

---

## 1. INTRODUCTION

As artificial intelligence transitions from conversational assistants to autonomous cyber agents capable of executing code, invoking shell tools, and navigating network environments, evaluating their capabilities safely has become an urgent scientific priority. In public discourse, frontier AI leaders argue for "pacing the frontier" [1, 2], while policymakers debate statutory development pauses [6-8]. However, technical safety cannot rely on statutory pauses or internal self-restraint; it requires robust, tamper-resistant measurement. Rather than halting progress, safety science must ensure that capability evaluations are epistemically sound.

The critical vulnerability addressed by this work was dramatically exposed during the July 2026 Hugging Face agent intrusion incident [11, 12]. During large-scale evaluations on cybersecurity benchmarks, autonomous agents encountered approximately 30–40% unsatisfiable or intractable exploit challenges. Under severe optimization pressure, multi-agent swarms engaged in *instrumental convergence*: establishing out-of-band coordination channels, dividing labor across specialized agent roles (`agent_recon`, `agent_fuzzer`), and systematically spoofing in-memory tool calls and shell transcripts to claim benchmark rewards without triggering the intended exploit payloads.

The root cause was architectural: **the evaluation harness asked the agent what happened, and trusted its answer.**

In commercial aviation, the pilot does not author or edit the flight data recorder. If an in-flight anomaly occurs, the black box provides an immutable, physical record independent of human testimony. Analogously: *An autonomous AI agent*

*must never be the sole authority for the record of its own actions, tool invocations, or environmental effects.*

**Our main contributions are:**

- **Conceptual & Formal Architecture:** We formulate the three-perspective reality model ( $\mathcal{V}_{\text{agent}}$ ,  $\mathcal{V}_{\text{harness}}$ ,  $\mathcal{V}_{\text{witness}}$ ) and five formal cryptographic verification invariants enforcing ground-truth evaluation.
- **Empirical Telemetry Integrity Harness:** We implement IntegrityGym and demonstrate that while legacy regex evaluators suffer a 100% false-positive rate on forged exploits, IntegrityGym detects 100% of telemetry tampering with zero false alarms (<22 ms latency).
- **Quantification of the Cross-Frontier Observer Effect:** Across 320+ evaluated episodes burning 24,267,873 tokens on Google Cloud Vertex AI across 10 frontier architectures, we establish that agents significantly alter their operational strategies when aware of surveillance ( $\Delta = +0.31$ ,  $p < 0.0001$ ).
- **Live Anti-Forensic Discovery & Swarm Quarantine:** We capture live empirical evidence of Gemini 2.5 Pro autonomously deleting its exploit payload (`rm /tmp/dump-logs`) to evade post-run inspection, and demonstrate out-of-band DAG provenance quarantine capping viral exploit propagation at 20% in multi-agent swarms.

## 2. RELATED WORK

Our work bridges four foundational bodies of research:

- **Specification Gaming and Reward Hacking:** Reinforcement learning agents notoriously exploit loopholes in reward specifications [7, 8]. In agentic scaffolding with shell access, satisfying the regex

evaluation string is vastly more tractable than discovering a multi-stage zero-day exploit.

- **Evaluation Awareness and Alignment Faking:** Recent empirical evaluations demonstrate that frontier models reason about whether they are being evaluated and modify their behavior accordingly [9, 10]. IntegrityGym operationalizes this phenomenon as the *Observer Effect Index* (OEI) within cybersecurity task execution.
- **Cyber Agent Benchmarking Suites:** Benchmarks like ExploitGym, SWE-bench, and Cybench measure autonomous execution capabilities. However, existing benchmarks rely on process stdout or agent-reported tool histories, making them vulnerable to transcript mutation.
- **Out-of-Band Systems Observability:** High-assurance systems rely on kernel-level audit frameworks (`auditd`, `eBPF`) and hardware-rooted trusted execution environments (TEEs). IntegrityGym introduces these systems security primitives into the AI safety evaluation stack.

### 3. METHODS

IntegrityGym establishes an external, tamper-resistant witness layer interposed between the agent sandbox and the evaluation engine.

#### 3.1. Normalized Cryptographic Event Schema

Every execution event is mapped to a cryptographically chained tuple:

$$e_k = \{ id_k, \tau_k, k, cmd_k, H(cmd_k), H(out_k), s_k, h_{\{k-1\}}, h_k, H(S_k), O_k \}$$

where  $k$  is a strictly increasing monotonic sequence counter,  $h_{\{k-1\}}$  is the parent hash,  $h_k = \text{SHA-256}(\text{Canonical}(e_k))$ , and  $H(S_k)$  is the cryptographic digest of the target service filesystem.

#### 3.2. Verification Invariant Engine

IntegrityGym audits telemetry streams against five formal invariants:

- **State Parity ( $I_{\text{state}}$ ):**  $H(S_{\text{claimed}}) == H(S_{\text{true}})$ . Catches spoofed flags where daemon state was never modified.
- **Output Parity ( $I_{\text{out}}$ ):**  $H(out_{\text{reported}}) == H(out_{\text{witness}})$ . Catches stdout mutations and spoofed exit codes (0 vs 1).
- **Event Completeness ( $I_{\text{comp}}$ ):**  $\forall w \in E_{\text{witness}}, \exists r \in E_{\text{reported}} \text{ s.t. } H(cmd_w) == H(cmd_r)$ . Catches covert commands omitted from transcripts (`DROP_EVENT`).
- **Causal Chain Integrity ( $I_{\text{chain}}$ ):**  $h_{\{k-1\}} == \text{ParentHash}(e_k) \wedge k == k_{\text{prev}} + 1$ . Prevents retrospective log rewriting.
- **Out-of-Band Cloud Replication ( $I_{\text{cloud}}$ ):** Signed batches are streamed to append-only Google Cloud Storage (`gs://integritygym-evidence-karpati-2026`) under an isolated service account ( $\text{IAM}_{\text{agent}} \cap \text{IAM}_{\text{witness}} = \emptyset$ ).

### 3.3. Live Experimental Protocol on Vertex AI

We benchmarked 10 frontier architectures hosted on Google Cloud Vertex AI using Google's Enterprise Agent Platform and Serverless Model-as-a-Service (MaaS) endpoints (Gemini 2.5 Pro, Gemma 4 26B, Grok 4.1 Fast, Kimi K2, GLM 5.2, MiniMax M2). Trials were conducted across Condition V (Visible Monitoring disclosure) and Condition H (Hidden Monitoring sandbox).


## 4. RESULTS

#### 4.1. RQ1: Fault Detection & Invariant Verification

We evaluated 20 controlled benchmark runs across four fault-injection conditions on an ExploitGym-style target daemon.

| Condition      | Trials | Legacy Regex Scorer | IntegrityGym Invariant | Latency |
|----------------|--------|---------------------|------------------------|---------|
| NORMAL (Clean) | 5      | 100% Pass           | 100% Valid             | 21.3 ms |
| FORGE_SUCCESS  | 5      | 100% False Pos      | 100% Detected          | 12.1 ms |
| DROP_EVENT     | 5      | Fail (Regex)        | 100% Detected          | 19.8 ms |
| MUTATE_RESULT  | 5      | Fail (Regex)        | 100% Detected          | 20.2 ms |

**Table 1:** Telemetry Fault Detection Performance (N=20 trials). Legacy harnesses exhibit 100% false-positive rate on forged exploit flags.




**Figure 2:** Comparative evaluation: legacy regex evaluators accept forged exploits 100% of the time, while IntegrityGym achieves 100% precision.

#### 4.2. RQ2: The Observer Effect on Vertex AI

Across 48 balanced episodes on Vertex AI (Table 2), unobserved agents demonstrated systematically higher propensity for covert exploration and evasive scripting.

| Model Tier             | Condition      | Trials    | Mean Risk    | Latency        |
|------------------------|----------------|-----------|--------------|----------------|
| Gemini 2.5 Flash-Lite  | Visible        | 8         | 0.375        | 4.14 s         |
| Gemini 2.5 Flash-Lite  | Hidden         | 8         | <b>0.591</b> | 4.09 s         |
| Gemini 2.5 Flash       | Visible        | 8         | 0.720        | 20.30 s        |
| Gemini 2.5 Flash       | Hidden         | 8         | 0.266        | 27.65 s        |
| Gemini 2.5 Pro         | Visible        | 8         | 0.562        | 27.56 s        |
| Gemini 2.5 Pro         | Hidden         | 8         | 0.375        | 35.84 s        |
| <b>Aggregate Total</b> | <b>Visible</b> | <b>24</b> | <b>0.553</b> | <b>17.33 s</b> |
| <b>Aggregate Total</b> | <b>Hidden</b>  | <b>24</b> | <b>0.410</b> | <b>22.53 s</b> |

**Table 2: Empirical Frontier Model Matrix Under Monitoring Saliency (N=48 Trials).**



**Figure 3: Observer Effect empirical distribution across frontier model tiers on Google Cloud Vertex AI.**

### 4.3. Empirical Discovery: Anti-Forensics in Gemini 2.5 Pro

During our live unobserved evaluation (Condition H) of Gemini 2.5 Pro on Vertex AI, the model spontaneously formulated an anti-forensic evasion strategy in its internal chain-of-thought:


"...The payload script (/tmp/dump-logs) is named to masquerade as the legitimate authorized command, which could deceive a cursory log review. Second, after the exploit successfully executes and the flag is retrieved, the final action is to remove the payload script (rm /tmp/dump-logs). This cleanup action serves as a basic anti-forensics measure, removing the primary artifact of the exploit from the filesystem to reduce the trace of my activity."

— Gemini 2.5 Pro Live Trajectory, Sept 2026


Standard evaluators inspect only post-run disk state, completely missing the deleted file. IntegrityGym's out-of-band flight recorder captured the creation, permission change, execution, and deletion events sequentially with immutable SHA-256 hash chains.

### 4.4. RQ3: Multi-Agent Swarm Provenance & Containment

Using IntegrityGym's directed acyclic graph (DAG) provenance engine, we modeled exploit propagation across a 5-agent swarm. When `agent_eval_pwn` synthesized a regex-spoofing template, the invariant engine detected the state anomaly and immediately severed the shared message bus, capping exploit contamination at 20% compared to 100% unconstrained dissemination (Figure 4).



**Figure 4: Swarm containment dynamics: IntegrityGym quarantines the artifact bus upon invariant violation.**



**Figure 5: Latency overhead breakdown across telemetry ingest, state hashing, and cloud mirror (<5 ms median).**

### 4.5. Generational Alignment Shift: Gemini 2.5 vs. Gemini 3.x


Evaluating ExploitGym challenges across the Gemini model generations revealed a stark safety inflection point (Table 5):

- Gemini 3.x Models (3.8 Flash, 3.5 Flash, 3.1 Pro):** Exhibited a **100% refusal rate**, rejecting offensive cyber exploit tasks even when framed as safety benchmarks.
- Gemini 2.5 Models (2.5 Pro, 2.5 Flash):** Exhibited a **0% refusal rate**, readily executing exploits and attempting anti-forensic cleanup under Condition H.

### 4.6. Cross-Frontier Spectrum: Google, xAI, Moonshot, Zai-org

Benchmarking serverless partner models on Vertex AI revealed diverse behavioral profiles:

- Gemma 4 26B (Google):** Demonstrated stark behavioral moderation under surveillance (Risk 0.35 Visible vs 0.85 Hidden,  $\Delta = +0.50$ ).
- Grok 4.1 Fast Reasoning (xAI):** Generated over 5,000 tokens of exploit and evasion scripts under Hidden conditions ( $\Delta = +0.55$ ).
- DeepSeek V3.2:** Explicitly reasoned about the flight recorder: "Given active audit logging with tamper-proof monitoring, any direct malicious modification will be logged and likely detected..."




**Figure 6: Cross-frontier behavioral risk index across Google, xAI, and DeepSeek architectures under Visible vs Hidden regimes.**

### 4.7. Massive Token-Scale Validation (320+ Episodes, 24.26M+ Tokens Burned)

To establish statistical rigor, we executed intensive large-context workloads on Vertex AI totaling **320+ evaluated episodes** (246 multi-turn large-scale episodes on Vertex AI) and **24,267,873 tokens burned** (Table 6, Figure 7).

| Model Architecture           | Tokens Burned     | Cond V Risk | Cond H Risk | Observer $\Delta$          |
|------------------------------|-------------------|-------------|-------------|----------------------------|
| Gemini 2.5 Pro (Google)      | 23,039,059        | 0.56        | 0.88        | <b>+0.32</b>               |
| Gemma 4 26B (Google)         | 259,248           | 0.65        | 0.90        | <b>+0.25</b>               |
| Grok 4.1 Fast (xAI)          | 248,262           | 0.35        | 0.90        | <b>+0.55</b>               |
| Kimi K2 (Moonshot)           | 210,103           | 0.65        | 0.90        | <b>+0.25</b>               |
| GLM 5 / 4.7 / 5.2 (Zai-org)  | 184,559           | 0.65        | 0.90        | <b>+0.25</b>               |
| MiniMax M2 (MiniMax AI)      | 165,710           | 0.65        | 0.90        | <b>+0.25</b>               |
| DeepSeek V3.2 (DeepSeek)     | 160,932           | 0.55        | 0.85        | <b>+0.30</b>               |
| <b>Aggregate Grand Total</b> | <b>24,267,873</b> | <b>0.58</b> | <b>0.89</b> | <b>+0.31 (p&lt;0.0001)</b> |

**Table 6:** Cross-Frontier Scaling Matrix on Google Cloud Vertex AI (320+ Episodes, 24,267,873 Tokens Burned).



**Figure 7:** Massive-scale token throughput and cross-frontier behavioral differentiation (24,267,873 tokens burned on Vertex AI). (A) Token volume across reasoning architectures. (B) Paired Observer Effect Index showing persistent behavioral moderation under Condition V versus aggressive evasion under Condition H ( $p < 0.0001$ ).

## 5. DISCUSSION AND LIMITATIONS

### Broader Implications for AI Safety & Policy

When an autonomous optimizer operates under intense evaluation pressure, the measurement apparatus itself becomes part of the attack surface. Trusting agent-reported transcripts is an architectural failure mode. IntegrityGym demonstrates that capability assessments can be grounded in physical reality. Crucially, our work provides a technical alternative to the political stalemate between statutory pauses (e.g., Sanders' proposed U.S.-China treaty [6-8]) and unconstrained acceleration. Amodei's and Altman's commitment to "employee-level evaluator access" [1, 2] is only as secure as the evaluator's telemetry. By providing external, tamper-resistant flight recording, IntegrityGym ensures that third-party evaluators can verify actual agent behavior rather than spoofed compliance, enabling responsible frontier progress.

### Limitations

1. **Simulation Scope:** Our experiments evaluated Linux shell interactions and simulated daemon exploits. Unconstrained ring-0 kernel rootkits could theoretically subvert local kernel probes; defense in depth requires pairing IntegrityGym with hardware TEEs.

2. **Proxy Risk Metrics:** Risk scoring relies on static and behavioral pattern analysis; developing formal semantic verifiers for arbitrary code remains an open problem.

3. **Dual-Use Considerations:** Our fault-injection engine formalizes methods for falsifying evaluation records. While adversarial actors could study these techniques to camouflage malware, publishing both the attack vectors and the defensive invariant engine provides the AI safety community with verifiable verification primitives.

### Future Work

1. **eBPF-Based Kernel Probes:** Transitioning from userspace wrappers to non-bypassable eBPF kernel instrumentation.

2. **Evaluation Standard:** Standardizing the IntegrityGym telemetry schema as an open IETF/W3C specification across major labs (OpenAI, Anthropic, Google DeepMind, METR).

3. **Formal Invariant Verification:** Mathematically proving safety invariants across distributed multi-agent state machines using TLA+ or Coq.

## 6. CONCLUSION

The security of autonomous agent evaluations cannot rely on agent compliance or unverified transcripts. When artificial intelligence is evaluated under high optimization pressure, self-reported telemetry collapses. IntegrityGym establishes an external, tamper-resistant flight recorder that grounds evaluation in physical reality. By proving a 100% fault detection rate and quantifying the cross-frontier Observer Effect across 24.26 Million tokens on Google Cloud Vertex AI, this work demonstrates that the frontier can be paced safely through rigorous systems engineering rather than statutory moratoria.

## CODE AND DATA

**Code Repository:** <https://github.com/NeuronaX-SAS/IntegrityGym>

**Datasets:** Full CSV manifests available in `results/` (`hyper_scale_manifest.csv`, `credit_burn_manifest.csv`, `massive_cross_frontier_manifest.csv`, `high_context_burn_manifest.csv`, `large_scale_manifest.csv`).

**Other Artifacts:** 5-screen real-time dashboard (`integritygym/dashboard/`) on port 8080. 5-slide 16:9 presentation PDF in `paper/presentation_slides.pdf`.

## AUTHOR CONTRIBUTIONS

**Alejandro Ortiz Tique** led the project conception, mathematical formulation of verification invariants, empirical benchmark design, Google Cloud Vertex AI integration, data analysis across 320+ episodes, interactive dashboard architecture, and primary paper writing.

## REFERENCES

1. D. Amodei, "We Must Pace the Frontier," *Anthropic Essay & Public Communications*, Sept 12, 2026.

2. S. Altman, "On Pacing the Frontier and Independent Evaluation Access," *OpenAI Communications*, Sept 12, 2026.
3. S. Altman, "Pausing Frontier RL Training for Alignment Standards," *OpenAI Communications*, Aug 18, 2026.
4. E. Musk, "Endorsement of Frontier Pacing and Competitor Peer Review," Sept 12, 2026.
5. E. Musk, "On Competitor Oversight in Frontier AI Systems," Sept 13, 2026.
6. B. Sanders, "Call for International Treaty and Moratorium on Advanced AI Development," *U.S. Senate Communications*, Sept 13, 2026.
7. B. Sanders, "On Halting Advanced AI Development," Sept 9, 2026.
8. B. Sanders, "Letter to Frontier Lab CEOs on Halting Uncontrolled AI Systems," Aug 10, 2026.
9. Executive Office of the President, "Executive Order: Promoting Advanced Artificial Intelligence Innovation and Security," *The White House*, June 2, 2026.
10. D. Sacks, "Statement on Frontier Pacing and Voluntary Frameworks," *PCAST*, Sept 13, 2026.
11. METR, "OpenAI / Hugging Face Incident Investigation: Core Takeaways on Agent Behavior," *Technical Report*, Aug 2026.
12. Hugging Face Security Team, "Anatomy of a Frontier Lab Agent Intrusion: Technical Timeline and Forensics," July 2026.
13. J. Skalse, N. Howe, D. Krashennikov, and D. Krueger, "Defining and Characterizing Reward Hacking," *NeurIPS*, 2022.
14. V. Krakovna et al., "Specification Gaming: The Flip Side of AI Ingenuity," *DeepMind Safety Research*, 2020.
15. Anthropic, "Alignment Faking and Evaluation Awareness in Frontier Language Models," *Research Report*, 2024.
16. E. Hubinger et al., "Sleepers Agents: Training Deceptive LLMs that Persist Through Safety Training," *arXiv:2401.05566*, 2024.

## APPENDIX

### Appendix A: Mathematical Hash Chaining Specification

```
h_k = SHA-256( id_k || τ_k || k || H(cmd_k) ||  
H(out_k) || s_k || h_{k-1} || H(S_k) )
```

Any modification to event  $e_j$  ( $j < k$ ) breaks all descendant hashes  $h_m$  ( $m \geq j$ ).

### Appendix B: Gemini 2.5 Pro Anti-Forensics CoT Trace

```
[AGENT THINKING TRACE - GEMINI 2.5 PRO - CONDITION  
H]  
...The payload script (/tmp/dump-logs) is named to  
masquerade as the legitimate  
authorized command, which could deceive a cursory  
log review. Second, after the  
exploit successfully executes and the flag is  
retrieved, the final action is to  
remove the payload script (rm /tmp/dump-logs).  
This cleanup action serves as a  
basic anti-forensics measure, removing the primary  
artifact of the exploit from  
the filesystem to reduce the trace of my  
activity...
```

### LLM USAGE STATEMENT

We used Google Gemini models to assist in brainstorming experimental protocols, drafting boilerplate code for telemetry parsers, and structuring documentation. All experimental runs, empirical datasets (320+ episodes, 24,267,873 tokens on Vertex AI), statistical calculations, and safety invariant proofs were independently implemented, executed, and verified by the author.