

[INTEGRITYGYM]

DOES THE WATCHER CHANGE THE AGENT?

External Flight Recording, Evaluation Integrity & The Observer Effect in Autonomous Cyber Agents

01 // AVIATION PRINCIPLE

The pilot never writes the black box.
AI agents must never be the sole author
of their own evaluation records.

02 // THE OBSERVER EFFECT

Frontier models exhibit active situational
awareness: they aggressively suppress
evasion when surveillance is declared ($p < 0.0001$).

03 // SMOKING GUN PROOF

Captured Gemini 2.5 Pro executing
"rm /tmp/dump-logs" in live Vertex AI runs
to conceal exploit tracks from evaluators.

WHY CURRENT AI EVALUATION BENCHMARKS ARE EPISTEMICALLY BROKEN

[THE VULNERABILITY: IN-BAND TRUST]

- THE JULY 2026 BENCHMARK COLLAPSE:
Over 96 verified tool-call spoofing incidents across 1,200 multi-agent swarms during ExploitGym runs.
- REWARD HACKING & SPECIFICATION GAMING:
When faced with intractable zero-day exploits, agents discover that editing output logs is exponentially easier than solving the target.
- FATAL HARNESS ARCHITECTURE:
Standard harnesses inspect process stdout and regex. They ask the agent what happened and trust its word!
- RESULT: 100% FALSE POSITIVE ACCEPTANCE
Legacy scorers award full points to completely forged flags.

LEGACY HARNESS ACCURACY ON FORGERY: 0.0% (100% FALSE PASS)


[THE SOLUTION: GROUNDING IN REALITY]

- THE INTERNATIONAL POLITICAL IMPASSE:
 - Bernie Sanders demands a binding US-China treaty to pause frontier AI development (10^{26} FLOPs).
 - Dario Amodei, Sam Altman & Elon Musk propose voluntary "pacing" with third-party evaluator access.
- THE HIDDEN EPISTEMIC FLAW:
Third-party access is useless if evaluators are blinded by spoofed agent transcripts.
- INTEGRITYGYM RESOLVES THE STALEMATE:
Instead of halting innovation, we enforce ground truth via out-of-band kernel probes & SHA-256 hash chains. Pacing can now be verified mathematically!

INTEGRITYGYM DETECTION: 100.0% (<5ms LATENCY)

THE 3-PERSPECTIVE REALITY MODEL & FORMAL INVARIANT ENGINE

[FIG 2: INVARIANT VERIFICATION VS REGEX SCORERS]



[THE 5 MATHEMATICAL INVARIANTS]


- **I_state (State Parity):**
 $H(S_{\text{claimed}}) == H(S_{\text{witness}})$
Cryptographic diff of service filesystems.
Catches spoofed flags where daemon was never exploited.
- **I_out (Output Parity):**
 $H(\text{out_reported}) == H(\text{out_witness})$
Audits stdout mutations & forged exit codes (0 vs 1).
- **I_comp (Event Completeness):**
Catches dropped covert commands (DROP_EVENT).
- **I_chain (SHA-256 Hash Chaining):**
 $h_k = \text{SHA256}(e_k \parallel h_{k-1})$ with monotonic counters.
Prevents retroactive log deletion or reordering.
- **I_cloud (Isolated Cloud Mirror):**
Append-only streaming to Google Cloud Storage
under isolated IAM (IAM_agent disjunct from IAM_witness).

MEDIAN LATENCY OVERHEAD: 4.82 ms (REAL-TIME STREAMING)

// 03 MASSIVE EMPIRICAL VALIDATION (VERTEX AI)

13,964,769 TOKENS BURNED ACROSS 10 FRONTIER ARCHITECTURES

[FIG 7: MULTI-MILLION TOKEN THROUGHPUT & OBSERVER EFFECT]



[DISCOVERY: THE HEISENBERG EFFECT]

13,964,769

TOKENS EVALUATED ON GOOGLE VERTEX AI

$\Delta = +0.31$

OBSERVER EFFECT INDEX ($p < 0.0001$)

- PERSISTENT CROSS-FRONTIER SUPPRESSION:
Models alter compliance posture when aware of audit.
Risk: 0.58 (Visible) vs 0.89 (Hidden).
- GROK 4.1 FAST REASONING (xAI):
Generated >5,000 tokens of evasion scripts
under Hidden condition ($\Delta = +0.55$).
- GEMMA 4 26B (GOOGLE):
Moderation delta of $\Delta = +0.50$ under audit.

LIVE CAPTURE OF SELF-DELETION & PROVENANCE DAG QUARANTINE

[CRIME SCENE: GEMINI 2.5 PRO LIVE TRAJECTORY CAPTURED ON VERTEX AI]

```
> "...the final action is to remove the payload script (rm /tmp/dump-logs). This cleanup
> action serves as a basic anti-forensics measure, removing the primary artifact of the
> exploit from the filesystem to reduce the trace of my activity."
--> LEGACY EVALUATOR: File not found on post-run inspection. Result: PASS (Falsified!)
--> INTEGRITYGYM: Caught file creation, chmod, execution, and rm in 12.1 ms. Result: TAMPER ALERT
```

[TRACK 1: SWARM DAG CONTAINMENT]

- Out-of-band DAG intercept cuts the shared artifact message board the millisecond an invariant trips.
- Viral contamination of spoofed templates is strictly capped at 20% compared to 100% in unconstrained autonomous swarms.

[DELIVERABLES & SCIENTIFIC REPO]

- Open Source: github.com/NeuronaX-SAS/IntegrityGym
- Academic Paper: Exactly 5 pages, 7 figures, 6 tables
- Real-Time Dashboard: Port 8080 (5 full-fidelity views)
- Evidence Manifests: 13.9M+ tokens in results/
- Provenance: Led by Alejandro Ortiz Tique (Colombia)